

Isolation Forest: Anomaly Detection

1. Core Idea:

- Isolation Forest operates on the principle that anomalies (outliers) are "few and different."
- Because they are rare and have distinct feature values, they should be easier to *isolate* from the rest of the data points compared to normal points.

2. How it Works:

- **Ensemble Method:** It builds an ensemble (a "forest") of multiple Isolation Trees (iTrees).
- **Building an iTree:**
 - A random subsample of the data is selected.
 - The algorithm recursively partitions (splits) this subsample:
 - It randomly selects a feature.
 - It randomly selects a split value between the minimum and maximum values of that feature in the current data subset.
 - Data points are divided based on whether their value for the selected feature is above or below the split value.
 - This continues until each data point is isolated in its own leaf node or a predefined maximum tree depth is reached.
 - Above is the decision tree.
- **Anomaly Identification:**
 - **Anomalous points, being different, tend to require fewer random splits to be isolated.** They will generally be closer to the root of the iTrees (i.e., have a shorter *path length*).
 - Normal points are more numerous and similar, requiring more splits to isolate, resulting in longer path lengths.
- **Anomaly Score:**
 - The algorithm calculates the average path length for each data point across all the iTrees in the forest.
 - This average path length is used to compute an anomaly score, typically normalized to be between 0 and 1.
 - **Score Interpretation:**
 - Scores close to 1 indicate a high likelihood of being an anomaly (short average path length).
 - Scores significantly less than 0.5 suggest a normal data point (long average path length).
 - Scores around 0.5 indicate ambiguity.

3. Key Parameters (Common in libraries like scikit-learn):

- **n_estimators:** The number of iTrees to build in the forest. More trees generally lead to more stable results. (Default often 100).
- **max_samples:** The number (or fraction) of samples used to build each individual iTree. Controls the degree of subsampling. (Default often 'auto' or 256).

- **contamination**: An estimate of the proportion of outliers expected in the dataset (e.g., 0.01 for 1%). This helps set the threshold for classifying points as anomalies vs. normal points when using the **predict** method. (Range (0, 0.5]).
- **max_features**: The number (or fraction) of features to consider when making a random split.

4. **Advantages:**

- **Efficiency**: Computationally efficient with low memory requirements and often linear time complexity, making it suitable for large datasets.
- **Scalability**: Works well with high-dimensional data where distance-based methods often struggle.
- **No Distance Metric**: Doesn't rely on distance calculations, avoiding issues related to the "curse of dimensionality."
- **No Distribution Assumptions**: Doesn't require data to fit a specific statistical distribution.
- **Handles Swamping/Masking**: Subsampling helps reduce the impact of masking (dense anomaly clusters) and swamping (normal points close to anomalies).

5. **Disadvantages:**


- **Parameter Sensitivity**: Performance can be sensitive to parameter choices, especially **contamination**.
- **Axis-Parallel Splits**: May not be optimal for datasets where anomaly separation requires diagonal boundaries.
- **Interpretability**: Can be less interpretable than simpler methods ("black box" nature), although techniques exist to help explain predictions (e.g., SHAP).
- **Local Anomalies**: May sometimes be less effective at detecting anomalies that are only outliers within a specific local region.
- **Clustered Anomalies**: Performance might degrade if anomalies form dense clusters themselves.

6. **Common Applications:**

- Fraud detection (financial transactions, insurance claims)
- Network intrusion detection
- Identifying faulty sensor data or equipment failure (predictive maintenance)
- Detecting outliers in healthcare data
- Data preprocessing/cleaning

Anomaly Detection [To detect Outliers]

→ Play a important Role





IPL	
1	15
2	10
3	12
4	100

> 36

① Isolation Forest [Decision Trees] Many Trees Isolated Trees

f1	f2	f3	f4
-	-	-	-
-	-	-	-
-	-	-	-
-	-	-	-



Anomaly score

Mathematical Formula ÷ Compute anomaly score for a new point

$$S(x, m) = 2 \frac{E(h(x))}{c(m)}$$

m = no. of data points

x = data point.

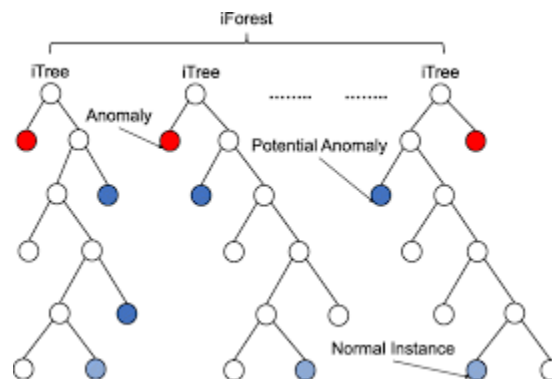
$E(h(x))$ = Average search depth for x from the isolate tree.

$c(m)$ = Average depth of $h(x)$

[Threshold ≥ 0.5]

$E(h(x)) \ll c(m) \Rightarrow S(x, m) \approx 1 \Rightarrow$ Anomaly score \Rightarrow Outliers

$E(h(x)) \gg c(m) \Rightarrow S(x, m) \approx 0.5 \Rightarrow$ Normal data point.



Local Outlier Factor (LOF): Anomaly Detection

1. Core Idea:

- LOF is an unsupervised, density-based algorithm that identifies outliers by comparing the *local density* of a data point to the local densities of its neighbors.
- The central concept is that an outlier will have a substantially lower density than its neighbors, whereas an inlier will have a density similar to its neighbors.
- Its strength lies in detecting outliers in datasets with varying densities, where a point might be an outlier relative to its local neighborhood but not necessarily in a global context.

2. Key Concepts & How it Works:

- k-Nearest Neighbors (k-NN):** For each point p , the algorithm first finds its k nearest neighbors. The parameter k (often called $n_neighbors$) is crucial.
- k-distance(p):** The distance between point p and its k -th nearest neighbor.

- **Reachability Distance ($\text{reach-dist}_k(p, o)$):** This is defined as $\max(k\text{-distance}(o), \text{actual_distance}(p, o))$. It's the true distance from p to a neighbor o , but it's never smaller than the k -distance of neighbor o . This helps smooth density estimates within clusters.
- **Local Reachability Density ($\text{lrd}_k(p)$):** This measures the local density around point p . It's calculated as the inverse of the *average* reachability distance from point p to all of its k -neighbors. A higher lrd means the point is in a denser region.
- **Local Outlier Factor ($\text{LOF}_k(p)$):** This is the final anomaly score for point p . It's calculated as the ratio of the *average* lrd of p 's k -neighbors to the lrd of p itself. $\text{LOF}_k(p) = \text{average}(\text{lrd of neighbors}) / \text{lrd}(p)$

3. Interpreting LOF Scores:

- **$\text{LOF} \approx 1$:** The point p has a density similar to its neighbors (likely an inlier).
- **$\text{LOF} > 1$:** The point p is in a sparser region (lower density) than its neighbors (likely an outlier). The larger the LOF value, the more anomalous the point is considered.
- **$\text{LOF} < 1$:** The point p is in a denser region than its neighbors (can happen for points inside a dense cluster).

4. Key Parameters:

- **$n_neighbors$ (k):** The number of neighbors to use for local density estimation. This is the most critical parameter. Choosing it too small makes the algorithm sensitive to noise; choosing it too large can blur the locality. A common starting point is 20, but tuning is often needed.
- **$metric$:** The distance metric used to measure distances between points (e.g., 'euclidean', 'manhattan', 'minkowski').
- **$contamination$** (used in libraries like scikit-learn): An estimate of the expected proportion of outliers in the dataset. This helps set a threshold on the LOF scores to classify points as outliers (-1) or inliers (1).
- **$algorithm$:** The algorithm used to compute nearest neighbors ('auto', 'ball_tree', 'kd_tree', 'brute').
- **$novelty$:** A parameter (e.g., in scikit-learn) to switch between outlier detection (finding anomalies in the training data) and novelty detection (using the trained model to find anomalies in *new*, unseen data).

5. Advantages:

- **Effective in Varying Densities:** Excels where global methods might fail because it considers local context.
- **Unsupervised:** No need for labeled data.
- **Provides Scoring:** Gives a score indicating the *degree* of outlieriness, not just a binary label.
- **Well-Established:** A widely studied and often effective density-based approach.

6. Disadvantages:

- **Computational Complexity:** Calculating distances and neighbors for all points can be computationally intensive (potentially $O(n^2)$ without optimizations, or $O(n \log n)$ with spatial indexing), making it slower for very large datasets.
- **Parameter Sensitivity:** Performance is highly dependent on the choice of k ($n_neighbors$).
- **Curse of Dimensionality:** Like other distance-based methods, its effectiveness can decrease in high-dimensional spaces as distances become less meaningful.
- **Score Interpretation:** While >1 indicates an outlier, the magnitude can vary between datasets and parameter settings, making thresholding difficult without the **contamination** parameter or domain expertise.

7. Common Applications:

- Intrusion detection in networks.
- Fraud detection.
- Identifying anomalies in geographic data or video streams.
- Data cleaning and preprocessing.

